

In [195]:

```
import tensorflow as tf
from tensorflow.keras import Sequential
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.layers import Conv2D, Activation, MaxPool2D, BatchNormalization, Flatten, Dense, Dropout
import pathlib
import scipy.io
import os, os.path, shutil
import pandas as pd
from PIL import Image
import numpy as np
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
```

In [208]:

```
def create_classes_folder(path_):
    images = [f for f in os.listdir(path_) if os.path.isfile(os.path.join(path_, f))]
    for image in images:
        folder_name = df[df.image == image]['class'].values[0]
        new_path = os.path.join(folder_path, folder_name)
        if not os.path.exists(new_path):
            os.makedirs(new_path)

        old_image_path = os.path.join(folder_path, image)
        new_image_path = os.path.join(new_path, image)
        shutil.move(old_image_path, new_image_path)
```

```
In [203]: def make_model(input_shape, num_classes, model_type_, aug_version_):
    inputs = keras.Input(shape=input_shape)
    # Image augmentation block
    if aug_version_ == 1:
        x = data_augmentation(inputs)
    elif aug_version_ == 2:
        x = aug(inputs)

    # Entry block
    x = layers.Rescaling(1.0 / 255)(x)
    if model_type_ == 1:
        x = layers.Conv2D(16, 3, padding='same', activation='relu')(x)
        x = layers.MaxPooling2D()(x)
        x = layers.Conv2D(32, 3, padding='same', activation='relu')(x)
        x = layers.MaxPooling2D()(x)
        x = layers.Conv2D(64, 3, padding='same', activation='relu')(x)
        x = layers.MaxPooling2D()(x)
        x = layers.Dropout(0.2)(x)
        x = layers.Flatten()(x)
        x = layers.Dense(128, activation='relu')(x)
    elif model_type_ == 2:
        x = layers.Conv2D(32, 3, strides=2, padding="same")(x)
        x = layers.BatchNormalization()(x)
        x = layers.Activation("relu")(x)

        x = layers.Conv2D(64, 3, padding="same")(x)
        x = layers.BatchNormalization()(x)
        x = layers.Activation("relu")(x)

    previous_block_activation = x # Set aside residual

    for size in [128, 256, 512, 728]:
        x = layers.Activation("relu")(x)
        x = layers.SeparableConv2D(size, 3, padding="same")(x)
        x = layers.BatchNormalization()(x)

        x = layers.Activation("relu")(x)
        x = layers.SeparableConv2D(size, 3, padding="same")(x)
        x = layers.BatchNormalization()(x)

        x = layers.MaxPooling2D(3, strides=2, padding="same")(x)
```

```
# Project residual
residual = layers.Conv2D(size, 1, strides=2, padding="same")(
    previous_block_activation
)
x = layers.add([x, residual])
previous_block_activation = x

x = layers.SeparableConv2D(1024, 3, padding="same")(x)
x = layers.BatchNormalization()(x)
x = layers.Activation("relu")(x)

x = layers.GlobalAveragePooling2D()(x)
x = layers.Dropout(0.5)(x)

activation = "softmax"
units = num_classes
outputs = layers.Dense(units, activation=activation)(x)
return keras.Model(inputs, outputs)
```

```
In [206]: def plot_metrics(history):
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']

    loss = history.history['loss']
    val_loss = history.history['val_loss']

    epochs_range = range(epochs)

    plt.figure(figsize=(8, 8))
    plt.subplot(1, 2, 1)
    plt.plot(epochs_range, acc, label='Training Accuracy')
    plt.plot(epochs_range, val_acc, label='Validation Accuracy')
    plt.legend(loc='lower right')
    plt.title('Training and Validation Accuracy')

    plt.subplot(1, 2, 2)
    plt.plot(epochs_range, loss, label='Training Loss')
    plt.plot(epochs_range, val_loss, label='Validation Loss')
    plt.legend(loc='upper right')
    plt.title('Training and Validation Loss')
    plt.show()
```

```
In [209]: def flatten_list(list_):
    flat_list = []
    for sublist in list_:
        for item in sublist:
            flat_list.append(item)
    return flat_list

def make_confusion_mtrx(model_, val_ds, save_path):
    y_pred = []
    y_true = []
    for x, y in val_ds.take(-1):
        y_pred.append(np.argmax(model_.predict(x), axis = 1))
        y_true.append(y)

    y_pred_fl = flatten_list(y_pred)
    y_true_fl = flatten_list(y_true)
    conf_mx = confusion_matrix(y_pred_fl, y_true_fl)
    df_cm = pd.DataFrame(conf_mx , index = [i for i in range(0,196)],
                          columns = [i for i in range(0,196)])
    cmap = sns.diverging_palette(230, 20, as_cmap=True)
    plt.figure(figsize = (100,100))
    cm_fig = sns.heatmap(df_cm, cmap=cmap, annot=True)
    fig = cm_fig.get_figure()
    fig.savefig(save_path)
    return cm_fig
```

```
In [210]: def make_a_prediction(model_, img_path):
    img = keras.preprocessing.image.load_img(
        img_path, target_size=(256,256))
    img_array = keras.preprocessing.image.img_to_array(img)
    img_array = tf.expand_dims(img_array, 0) # Create batch axis

    predictions = model_.predict(img_array)
    score = predictions[0]
    image = Image.open(img_path)
    display(image)
    return (f'The car on the image is most likely - {train_ds.class_names[np.argmax(score)]}')
```

```
In [5]: tf.__version__
```

```
Out[5]: '2.6.0'
```

```
In [2]: physical_devices = tf.config.list_physical_devices('GPU')
tf.config.experimental.set_memory_growth(physical_devices[0], True)

for device in physical_devices:
    tf.config.experimental.set_memory_growth(device, True)
```

```
In [49]: # uncomment for downloading data
# data_dir = tf.keras.utils.get_file(
#     "stanford_cars",
#     "http://ai.stanford.edu/~jkrause/car196/car_ims.tgz",
#     untar=True, extract= True)
# data_dir = pathlib.Path(data_dir)
```

```
In [5]: #data_dir = pathlib.Path(data_dir)
data_dir = 'C:/Users/kaara/.keras/datasets/car_ims'
image_count = len(list(data_dir.glob('/*/*.jpg')))
mat = scipy.io.loadmat('C:/Users/kaara/.keras/datasets/cars_annos.mat')
print(image_count)
```

```
16103
```

```
In [7]: data = []
for i in range (0,mat['annotations'].shape[1]):
    obj = mat['annotations'][0][i]
    im_path = obj[0][0].split('/')[1]
    class_name = mat['class_names'][0][obj[5][0][0]-1][0]
    data.append([im_path,class_name])
df = pd.DataFrame(data, columns = ['image','class'])
```

```
In [69]: # Uncomment when new data has been downloaded
# This step is needed because in archive all data store in one folder,
# to use tf.image_dataset_from_directory we need folder for each class

#create_classes_folder('C:/Users/kaara/.keras/datasets/car_ims')
```

```
In [70]: image_size = (256,256)
batch_size = 16

train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="training",
    seed=1337,
    image_size=image_size,
    batch_size=batch_size,
)
val_ds = tf.keras.preprocessing.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="validation",
    seed=1337,
    image_size=image_size,
    batch_size=batch_size,
)
```

```
Found 16185 files belonging to 196 classes.
Using 12948 files for training.
Found 16185 files belonging to 196 classes.
Using 3237 files for validation.
```

In [196]:

```
plt.figure(figsize=(16, 16))
for images, labels in train_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(train_ds.class_names[int(labels[i])]))
        plt.axis("off")
```

Tesla Model S Sedan 2012



Chevrolet Corvette Ron Fellows Edition Z06 2007



Lamborghini Diablo Coupe 2001



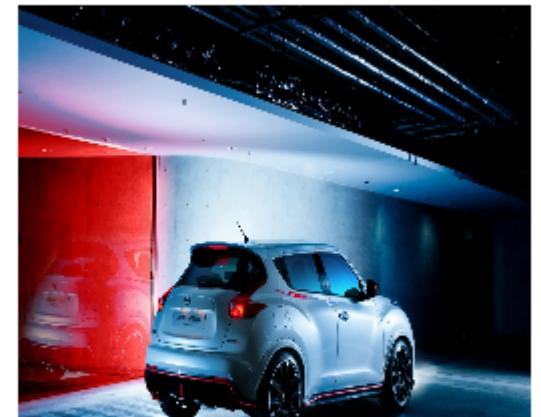
Hyundai Genesis Sedan 2012



Chevrolet Silverado 1500 Classic Extended Cab 2007



Nissan Juke Hatchback 2012





Chrysler Town and Country Minivan 2012



Tesla Model S Sedan 2012

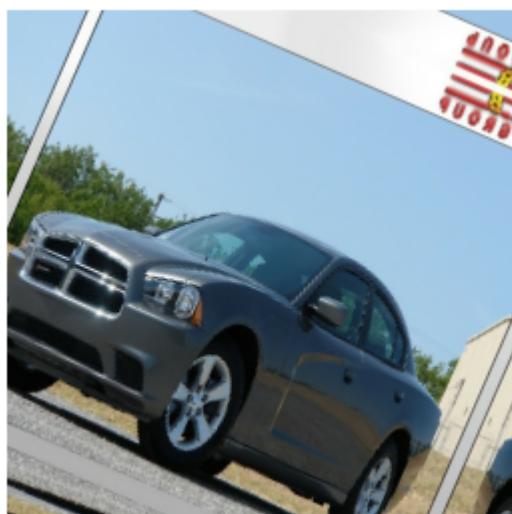
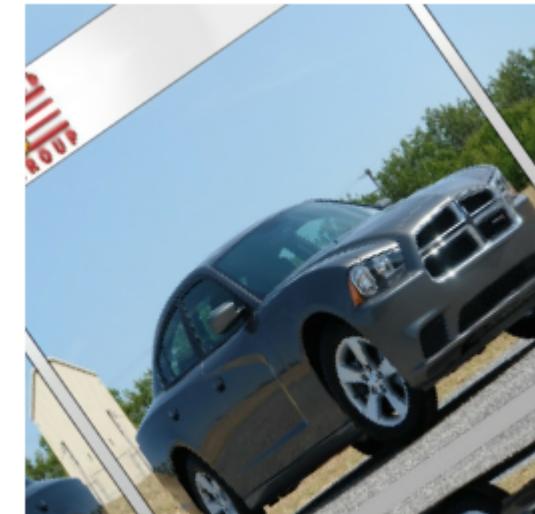
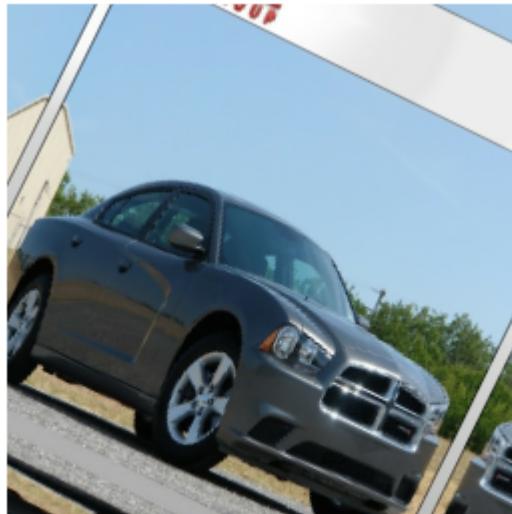


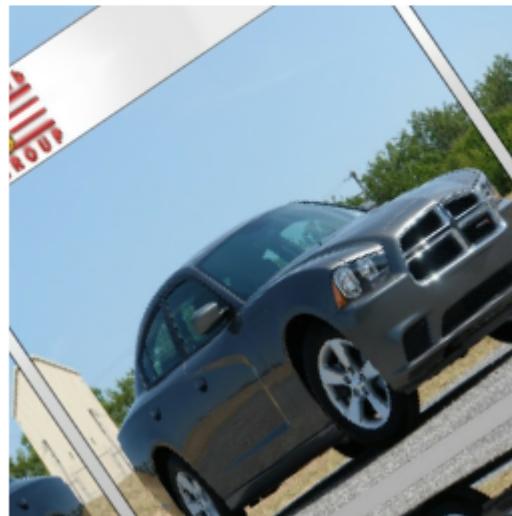
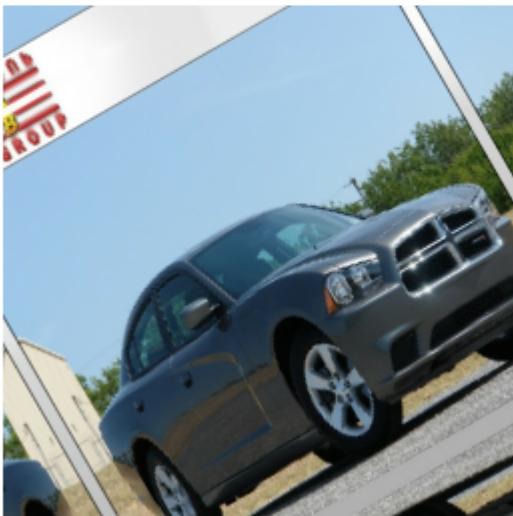
Bentley Arnage Sedan 2009



```
In [198]: data_augmentation = keras.Sequential(  
    [  
        layers.RandomFlip("horizontal"),  
        layers.RandomRotation(0.1),  
    ]  
)
```

```
In [199]: plt.figure(figsize=(16, 16))
for images, _ in train_ds.take(1):
    for i in range(9):
        augmented_images = data_augmentation(images)
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(augmented_images[0].numpy().astype("uint8"))
        plt.axis("off")
```





```
In [20]: AUTOTUNE = tf.data.AUTOTUNE  
  
train_ds = train_ds.cache().prefetch(buffer_size=AUTOTUNE)  
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

```
In [204]: model = make_model(input_shape=image_size + (3,), num_classes=196, model_type_= 1, aug_version_= 1)
model.summary()
max_pooling2d_4 (MaxPooling2 (None, 128, 128, 32) 0
conv2d_7 (Conv2D) (None, 128, 128, 32) 4640
max_pooling2d_5 (MaxPooling2 (None, 64, 64, 32) 0
conv2d_8 (Conv2D) (None, 64, 64, 64) 18496
max_pooling2d_6 (MaxPooling2 (None, 32, 32, 64) 0
dropout_1 (Dropout) (None, 32, 32, 64) 0
flatten (Flatten) (None, 65536) 0
dense_1 (Dense) (None, 128) 8388736
dense_2 (Dense) (None, 196) 25284
=====
Total params: 8,437,604
Trainable params: 8,437,604
```

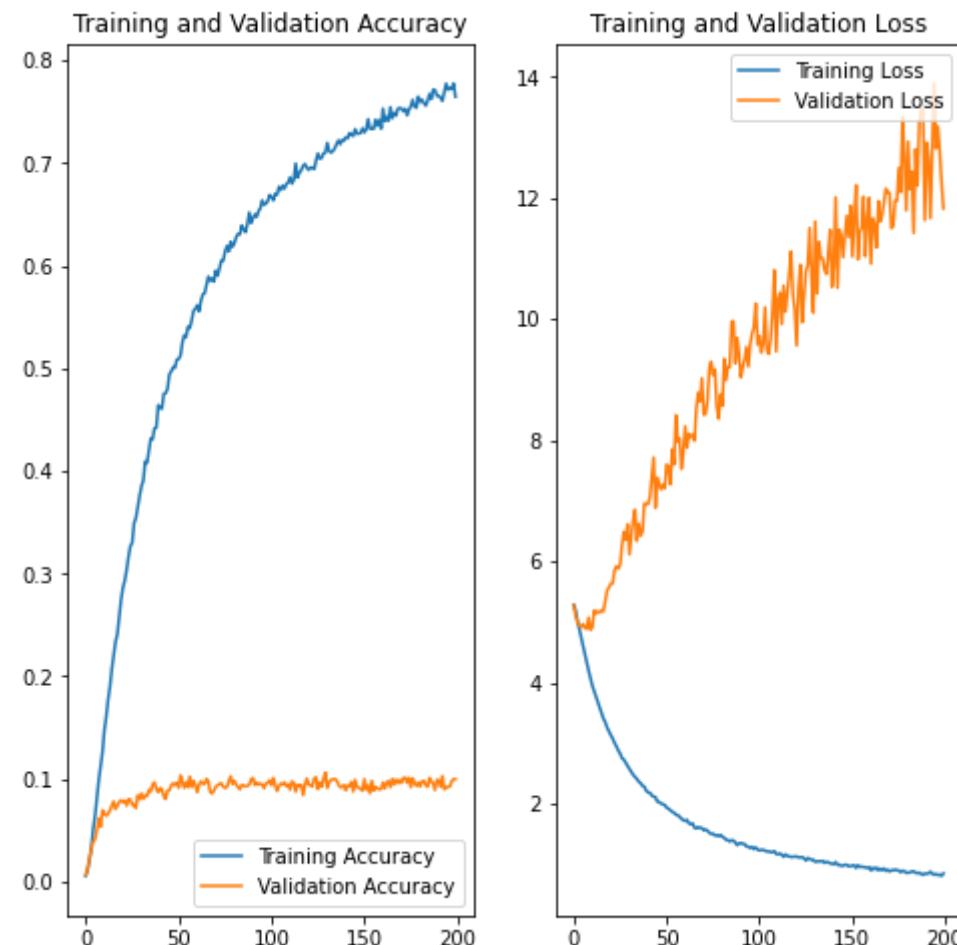
```
In [2]: print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
```

```
Num GPUs Available: 1
```

```
In [55]: epochs = 200
```

```
model.compile(  
    optimizer=keras.optimizers.Adam(1e-3),  
    loss=tf.keras.losses.SparseCategoricalCrossentropy(),  
    metrics=["accuracy"],  
)  
history_smp = model.fit(  
    train_ds, epochs=epochs, validation_data=val_ds,  
)  
810/810 [=====] - 19s 24ms/step - loss: 1.0913 - accuracy: 0.7049 - val_loss: 11.5021 - val_  
accuracy: 0.0924  
Epoch 129/200  
810/810 [=====] - 19s 23ms/step - loss: 1.0724 - accuracy: 0.7093 - val_loss: 10.7139 - val_  
accuracy: 0.1001  
Epoch 130/200  
810/810 [=====] - 19s 23ms/step - loss: 1.0761 - accuracy: 0.7100 - val_loss: 10.1031 - val_  
accuracy: 0.1063  
Epoch 131/200  
810/810 [=====] - 19s 23ms/step - loss: 1.0411 - accuracy: 0.7195 - val_loss: 11.6080 - val_  
accuracy: 0.0914  
Epoch 132/200  
810/810 [=====] - 19s 23ms/step - loss: 1.0470 - accuracy: 0.7117 - val_loss: 10.4184 - val_  
accuracy: 0.0964  
Epoch 133/200  
810/810 [=====] - 19s 23ms/step - loss: 1.0541 - accuracy: 0.7101 - val_loss: 11.2696 - val_  
accuracy: 0.0992  
Epoch 134/200  
810/810 [=====] - 19s 23ms/step - loss: 1.0474 - accuracy: 0.7115 - val_loss: 11.0040 - val_  
accuracy: 0.0995
```

```
In [56]: plot_metrics(history_smp)
```



That was just a simple CNN, which, actualy didn't succed on the validation data. So let's build more complex one.

```
In [57]: model_res = make_model(input_shape=image_size + (3,), num_classes=196, model_type_= 2, aug_version_= 1)
model_res.summary()
```

max_pooling2d_42 (MaxPooling2D)	(None, 32, 32, 256)	0	batch_normalization_60[0][0]
conv2d_54 (Conv2D)	(None, 32, 32, 256)	33024	add_20[0][0]
add_21 (Add)	(None, 32, 32, 256)	0	max_pooling2d_42[0][0] conv2d_54[0][0]
activation_61 (Activation)	(None, 32, 32, 256)	0	add_21[0][0]
separable_conv2d_49 (SeparableC	(None, 32, 32, 512)	133888	activation_61[0][0]
batch_normalization_61 (BatchNo	(None, 32, 32, 512)	2048	separable_conv2d_49[0][0]
activation_62 (Activation)	(None, 32, 32, 512)	0	batch_normalization_61[0][0]
separable_conv2d_50 (SeparableC	(None, 32, 32, 512)	267264	activation_62[0][0]
batch_normalization_62 (BatchNo	(None, 32, 32, 512)	2048	separable_conv2d_50[0][0]

Initialy, I've ran it on 200 epochs, and it has taken ~10 hours to train. In case, if something goes wrong with cells, results can be found in images, which I've put into the archive.

Results: loss: 0.0468 - accuracy: 0.9852 - val_loss: 2.0041 - val_accuracy: 0.7322

In [59]: epochs = 200

```
model_res.compile(  
    optimizer=keras.optimizers.Adam(1e-3),  
    loss=tf.keras.losses.SparseCategoricalCrossentropy(),  
    metrics=["accuracy"],  
)  
history = model_res.fit(  
    train_ds, epochs=epochs, validation_data=val_ds,  
)  
_accuracy: 0.7646  
Epoch 142/200  
810/810 [=====] - 165s 204ms/step - loss: 0.0643 - accuracy: 0.9795 - val_loss: 1.9566 - val_accuracy: 0.7152  
Epoch 143/200  
810/810 [=====] - 164s 203ms/step - loss: 0.0675 - accuracy: 0.9788 - val_loss: 1.7729 - val_accuracy: 0.7223  
Epoch 144/200  
810/810 [=====] - 165s 204ms/step - loss: 0.0788 - accuracy: 0.9761 - val_loss: 1.6163 - val_accuracy: 0.7377  
Epoch 145/200  
810/810 [=====] - 164s 203ms/step - loss: 0.0649 - accuracy: 0.9801 - val_loss: 1.8503 - val_accuracy: 0.7303  
Epoch 146/200  
810/810 [=====] - 163s 202ms/step - loss: 0.0577 - accuracy: 0.9824 - val_loss: 1.4636 - val_accuracy: 0.7640  
Epoch 147/200  
810/810 [=====] - 167s 206ms/step - loss: 0.0613 - accuracy: 0.9798 - val_loss: 1.7593 - val_accuracy: 0.7257  
Epoch 148/200
```

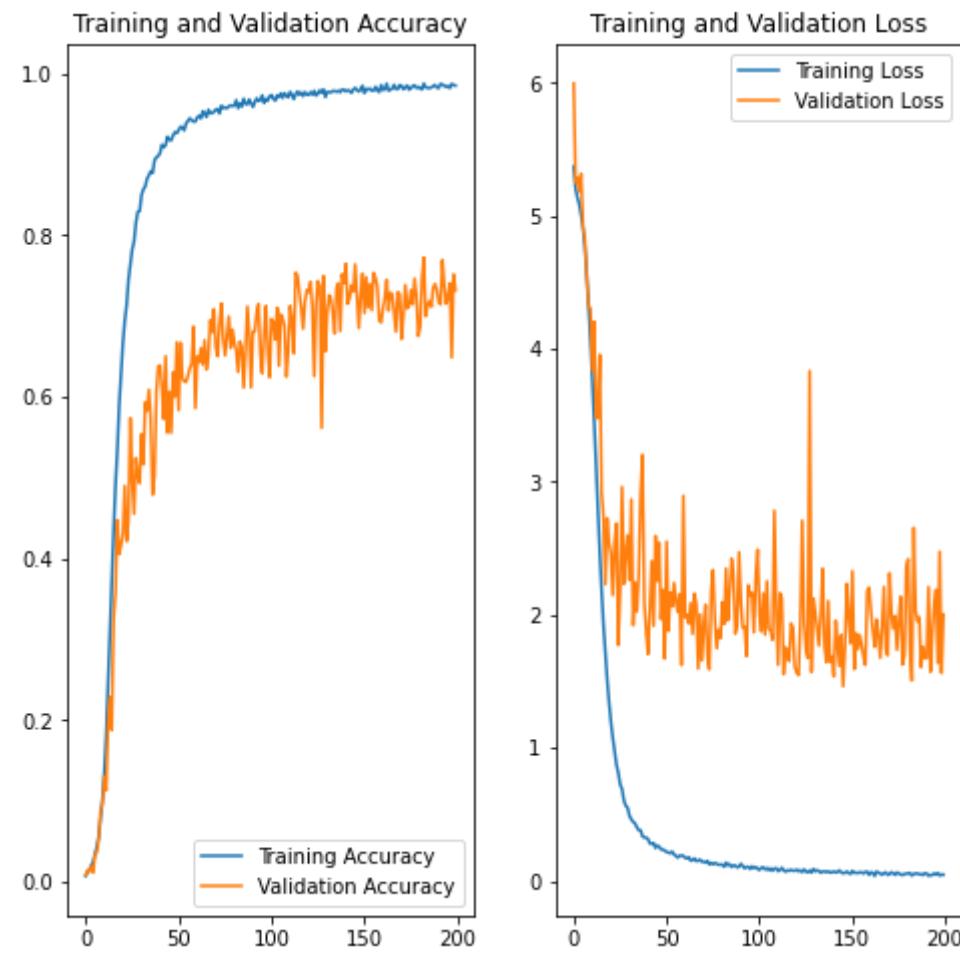
```
In [60]: acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

loss = history.history['loss']
val_loss = history.history['val_loss']

epochs_range = range(epochs)

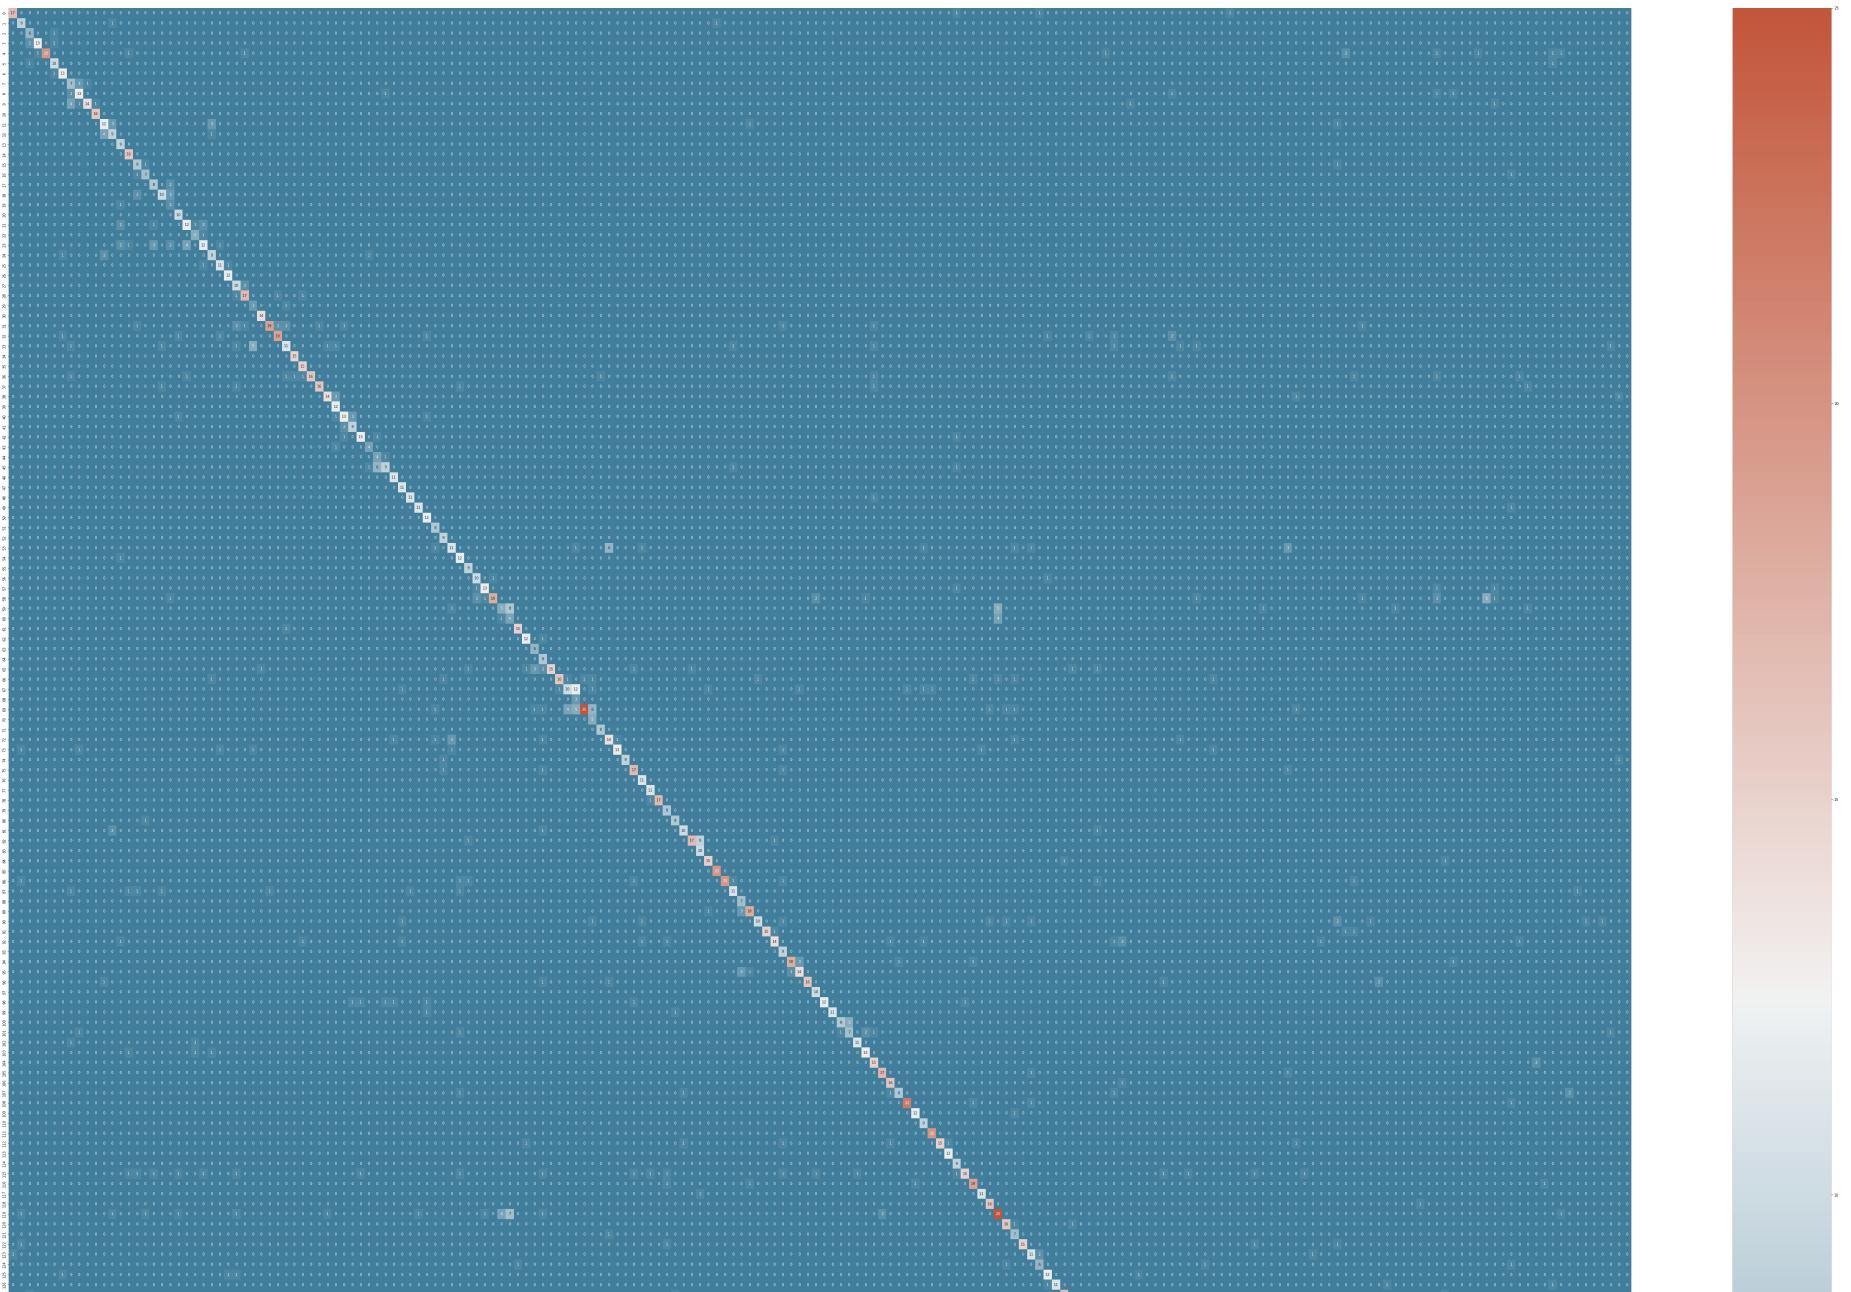
plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

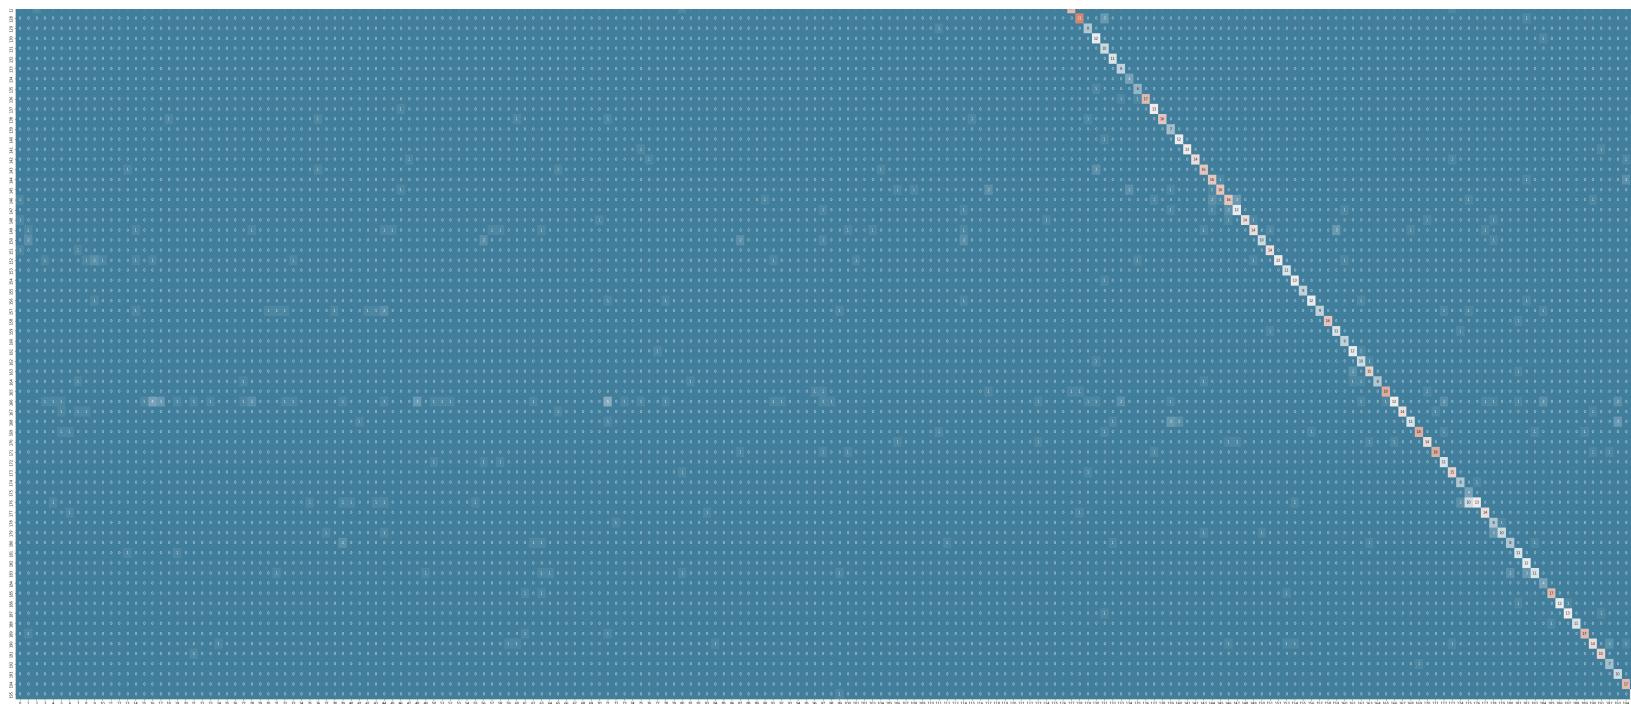
plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



```
In [81]: make_confusion_mtrix(model_ = model_res, validation_df = val_ds, save_path = 'confusion_matrix.jpg')
```

```
Out[81]: <AxesSubplot:>
```





Hint: click on the matrix, it is much bigger, than Jupier can display in the cell

```
In [82]: fig = cm_fig.get_figure()
fig.savefig('confusion_matrix.png')
```

I've noticed that the last 50 epochs hasn't made a great influence on overall accuracy, so I've retrained the model on 150 epochs and add slightly more complex data augmentation

```
In [10]: aug = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.5,0.3)
    ]
)
```

In [205]:

```
model_res_new = make_model(input_shape=image_size + (3,), num_classes=196, model_type_= 2, aug_version_= 2)
model_res_new.summary()
```

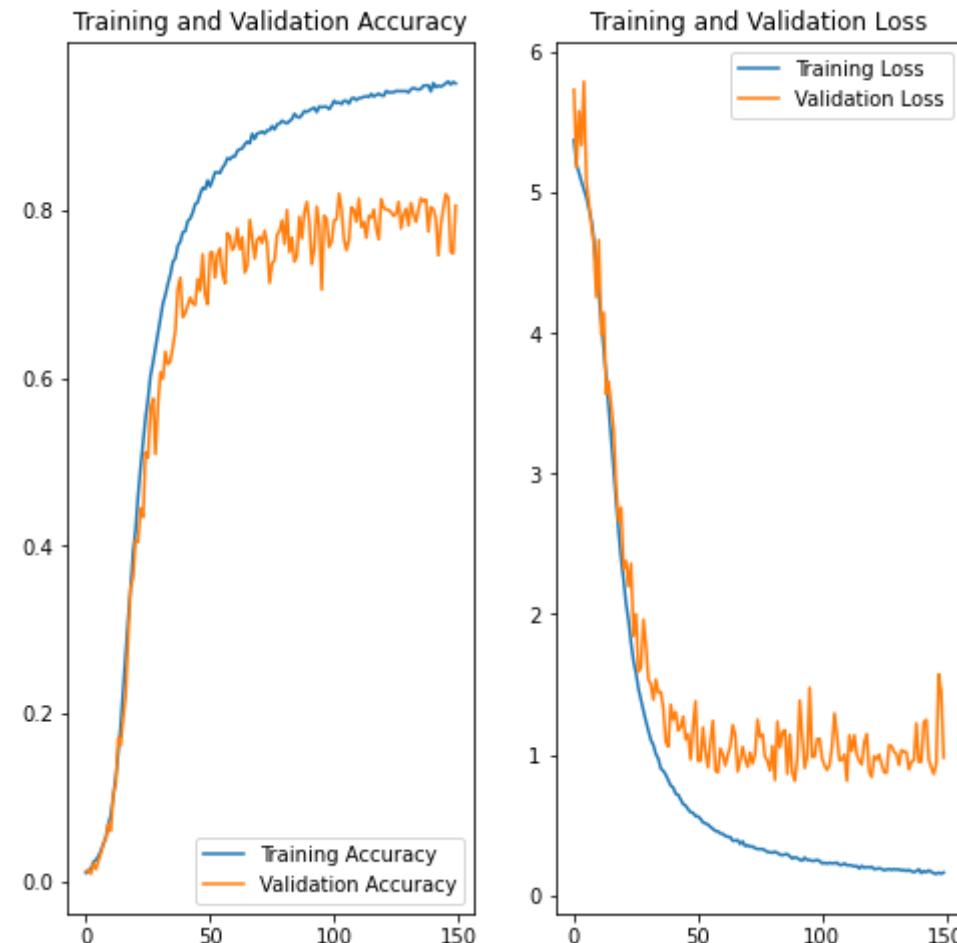
conv2d_13 (Conv2D)	(None, 16, 16, 512)	131584	add_5[0][0]
add_6 (Add)	(None, 16, 16, 512)	0	max_pooling2d_9[0][0] conv2d_13[0][0]
activation_19 (Activation)	(None, 16, 16, 512)	0	add_6[0][0]
separable_conv2d_15 (SeparableC	(None, 16, 16, 728)	378072	activation_19[0][0]
batch_normalization_19 (BatchNo	(None, 16, 16, 728)	2912	separable_conv2d_15[0][0]
activation_20 (Activation)	(None, 16, 16, 728)	0	batch_normalization_19[0][0]
separable_conv2d_16 (SeparableC	(None, 16, 16, 728)	537264	activation_20[0][0]
batch_normalization_20 (BatchNo	(None, 16, 16, 728)	2912	separable_conv2d_16[0][0]
max_pooling2d_10 (MaxPooling2D)	(None, 8, 8, 728)	0	batch_normalization_20[0][0]

In [13]: epochs = 150

```
callbacks = [
    keras.callbacks.ModelCheckpoint("save_at_{epoch}.h5"),
]

model_res_new.compile(
    optimizer=keras.optimizers.Adam(1e-3),
    loss=tf.keras.losses.SparseCategoricalCrossentropy(),
    metrics=["accuracy"],
)
history_new_res = model_res_new.fit(
    train_ds, epochs=epochs, callbacks=callbacks, validation_data=val_ds,
)
_accuracy: 0.7748
Epoch 80/150
810/810 [=====] - 169s 209ms/step - loss: 0.3048 - accuracy: 0.9061 - val_loss: 0.8870 - val_accuracy: 0.7875
Epoch 81/150
810/810 [=====] - 175s 216ms/step - loss: 0.3023 - accuracy: 0.9043 - val_loss: 1.0626 - val_accuracy: 0.7593
Epoch 82/150
810/810 [=====] - 170s 210ms/step - loss: 0.3089 - accuracy: 0.9030 - val_loss: 0.8206 - val_accuracy: 0.7998
Epoch 83/150
810/810 [=====] - 167s 206ms/step - loss: 0.3007 - accuracy: 0.9057 - val_loss: 1.2359 - val_accuracy: 0.7495
Epoch 84/150
810/810 [=====] - 167s 206ms/step - loss: 0.2895 - accuracy: 0.9079 - val_loss: 1.0517 - val_accuracy: 0.7671
Epoch 85/150
810/810 [=====] - 167s 206ms/step - loss: 0.2843 - accuracy: 0.9144 - val_loss: 1.1604 - val_accuracy: 0.7383
Epoch 86/150
```

```
In [207]: plot_metrics(history_new_res)
```



And now we can see, that adding just one more step into the data augmentation makes it ~5% more accurate and greatly reduce the validation loss even on lower numbers of epochs (150 vs 200 with less augmentation). And it is really great! It is much easier to preprocess images, than try to optimize the number and order of layers. and more time effective than trying to fit NN more.

In [39]: `make_a_prediction(model_res_new, "h2.jpg")`



Out[39]: 'The car on the image is most likely - HUMMER H2 SUT Crew Cab 2009'

Comment: perfect match!

```
In [40]: make_a_prediction(model_res_new,"audi_q3.jpg")
```



Out[40]: 'The car on the image is most likely - Audi R8 Coupe 2012'

Comment: So, it understand that it is an Audi, however it is Q3 Sportback while R8 is a sport car. Not very accurate, I supposed there had been some Audi crossovers in the dataset

```
In [41]: make_a_prediction(model_res_new, "Ferrari_SF90.jpg")
```



Out[41]: 'The car on the image is most likely - Ferrari 458 Italia Convertible 2012'

Comment: So, again, it is not 458 Italia, it is SF90, however, they are really quite similar

```
In [190]: make_confusion_mtrix(model_res_new, val_ds, 'conf_mtx.jpg')
```

```
Out[190]: <AxesSubplot:>
```

